

## Predicting House Prices Using Linear Regression

```
# 📦 Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
```

Methods for scaling, centering, normalization, binarization, and more.

```
# 📁 Load the dataset
housing = fetch_california_housing(as_frame=True)
data = housing.frame
```

```
# 🔎 Explore the data
print(data.info())
print(data.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
0   MedInc       20640 non-null   float64 
1   HouseAge     20640 non-null   float64 
2   AveRooms     20640 non-null   float64 
3   AveBedrms    20640 non-null   float64 
4   Population   20640 non-null   float64 
5   AveOccup     20640 non-null   float64 
6   Latitude     20640 non-null   float64 
7   Longitude    20640 non-null   float64 
8   MedHouseVal  20640 non-null   float64 
dtypes: float64(9)
memory usage: 1.4 MB
None
```


	MedInc	HouseAge	AveRooms	AveBedrms	Population	
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	
mean	3.870671	28.639486	5.429000	1.096675	1425.476744	
std	1.899822	12.585558	2.474173	0.473911	1132.462122	
min	0.499900	1.000000	0.846154	0.333333	3.000000	
25%	2.563400	18.000000	4.440716	1.006079	787.000000	
50%	3.534800	29.000000	5.229129	1.048780	1166.000000	
75%	4.743250	37.000000	6.052381	1.099526	1725.000000	
max	15.000100	52.000000	141.909091	34.066667	35682.000000	


	AveOccup	Latitude	Longitude	MedHouseVal
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.070655	35.631861	-119.569704	2.068558
std	10.386050	2.135952	2.003532	1.153956
min	0.692308	32.540000	-124.350000	0.149990
25%	2.429741	33.930000	-121.800000	1.196000
50%	2.818116	34.260000	-118.490000	1.797000
75%	3.282261	37.710000	-118.010000	2.647250
max	1243.333333	41.950000	-114.310000	5.000010

```
# Plot histogram of target
plt.hist(data['MedHouseVal'], bins=30, edgecolor='k')
plt.title("Distribution of Median House Values")
plt.xlabel("Median House Value ($100,000s)")
plt.ylabel("Frequency")
plt.show()
```

### Distribution of Median House Values



```
# Correlation matrix
plt.figure(figsize=(12, 8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
# ✎ Data preparation
X = data.drop('MedHouseVal', axis=1)
y = data['MedHouseVal']
```

```
# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)
```

```
# 🌈 Feature selection (based on correlation heatmap)
selected_features = ['MedInc', 'AveRooms', 'HouseAge']
X_selected = X_scaled[selected_features]

# ✎ Split the data
X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.2, random_state=42)
```

```
# 🖥 Train the model
model = LinearRegression()
model.fit(X_train, y_train)

▼ LinearRegression ⓘ ⓘ
LinearRegression() (module) preprocessing
```

```
Methods for scaling, centering, normalization, binarization, and more.


# 🕵️ Predict
y_pred = model.predict(X_test)
```

```
# 📈 Evaluate the model
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

print(f"R-squared: {r2:.4f}")
print(f"Mean Squared Error: {mse:.4f}")
```

```
R-squared: 0.4972
Mean Squared Error: 0.6589
```

```
# 📊 Visualize predictions
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual Median House Value")
plt.ylabel("Predicted Median House Value")
plt.title("Actual vs. Predicted House Values")
plt.grid(True)
plt.show()
```



Start coding or [generate](#) with AI.